

Kubernetes

Containers

- ▶ A good way to bundle and run our applications.
- ▶ In a production environment
 - ▶ We need to manage the containers that
 - ▶ Run the applications and ensure that there is no downtime.
 - ▶ If a container goes down, another container needs to start.
- ▶ Kubernetes
 - ▶ Provides us with a framework to run distributed systems resiliently.
 - ▶ It takes care of scaling and failover for our application, provides deployment patterns, and more.
 - ▶ Eg: kubernetes can easily manage a canary deployment for our system.

Kubernetes

- ▶ A portable, extensible, open source platform for managing
 - ▶ containerized workloads and services,
 - ▶ that facilitates both declarative configuration and automation.
- ▶ Has a large, rapidly growing ecosystem.
- ▶ Kubernetes services, support, and tools - widely available.

Kubernetes provides us with:

- ▶ Service discovery and load balancing
 - ▶ Kubernetes can expose a container using the DNS name or using their own IP address. If traffic to a container is high, Kubernetes is able to load balance and distribute the network traffic so that the deployment is stable.
- ▶ Storage orchestration
 - ▶ Kubernetes allows us to automatically mount a storage system of our choice, such as local storages, public cloud providers.

Kubernetes provides us with:

- ▶ Automated rollouts and rollbacks
 - ▶ We can describe the desired state for our deployed containers using Kubernetes
 - ▶ and it can change the actual state to the desired state at a controlled rate.
 - ▶ Eg: we can automate Kubernetes to create new containers for our deployment, remove existing containers and adopt all their resources to the new container.
- ▶ Automatic bin packing
 - ▶ We provide Kubernetes with a cluster of nodes that it can use to run containerized tasks.
 - ▶ We tell Kubernetes how much CPU and memory (RAM) each container needs.
 - ▶ Kubernetes can fit containers onto our nodes to make the best use of our resources.

Kubernetes provides us with:

▶ Self-healing

▶ Kubernetes

- ▶ restarts containers that fail,
- ▶ replaces containers,
- ▶ kills containers that don't respond to our user-defined health check, and doesn't advertise them to clients until they are ready to serve.

Kubernetes provides us with:

- ▶ Secret and configuration management
 - ▶ Kubernetes lets us store and manage sensitive information, such as passwords, OAuth tokens, and SSH keys.
 - ▶ we can deploy and update secrets and application configuration without rebuilding our container images, and without exposing secrets in our stack configuration.
- ▶ Batch execution
 - ▶ In addition to services, Kubernetes can manage our batch and CI workloads, replacing containers that fail, if desired.

Kubernetes provides us with:

- ▶ Horizontal scaling
 - ▶ Scale our application up and down with a simple command, with a UI, or automatically based on CPU usage.
- ▶ IPv4/IPv6 dual-stack
 - ▶ Allocation of IPv4 and IPv6 addresses to Pods and Services
- ▶ Designed for extensibility
 - ▶ Add features to our Kubernetes cluster without changing upstream source code.

Kubernetes provides us with:

- ▶ Secret and configuration management
 - ▶ Kubernetes lets us store and manage sensitive information, such as passwords, OAuth tokens, and SSH keys.
 - ▶ we can deploy and update secrets and application configuration without rebuilding our container images, and without exposing secrets in our stack configuration.
- ▶ Batch execution
 - ▶ In addition to services, Kubernetes can manage our batch and CI workloads, replacing containers that fail, if desired.

Kubernetes

- ▶ Not a traditional, all-inclusive PaaS system.
- ▶ operates at the container level rather than at the hardware level,
 - ▶ it provides some generally applicable features common to PaaS offerings, such as
 - ▶ deployment, scaling, load balancing, and
 - ▶ lets users integrate their logging, monitoring, and alerting solutions.
- ▶ It is not monolithic, and these default solutions are optional and pluggable.
- ▶ provides the building blocks for building developer platforms, but preserves user choice and flexibility where it is important.

Kubernetes:

- ▶ Does not limit the types of applications supported.
 - ▶ aims to support an extremely diverse variety of workloads,
 - ▶ including stateless, stateful, and data-processing workloads.
 - ▶ If an application can run in a container, it should run great on Kubernetes.
- ▶ Does not deploy source code and does not build our application.
 - ▶ CI/CD workflows are determined by organization cultures and preferences as well as technical requirements.

Kubernetes:

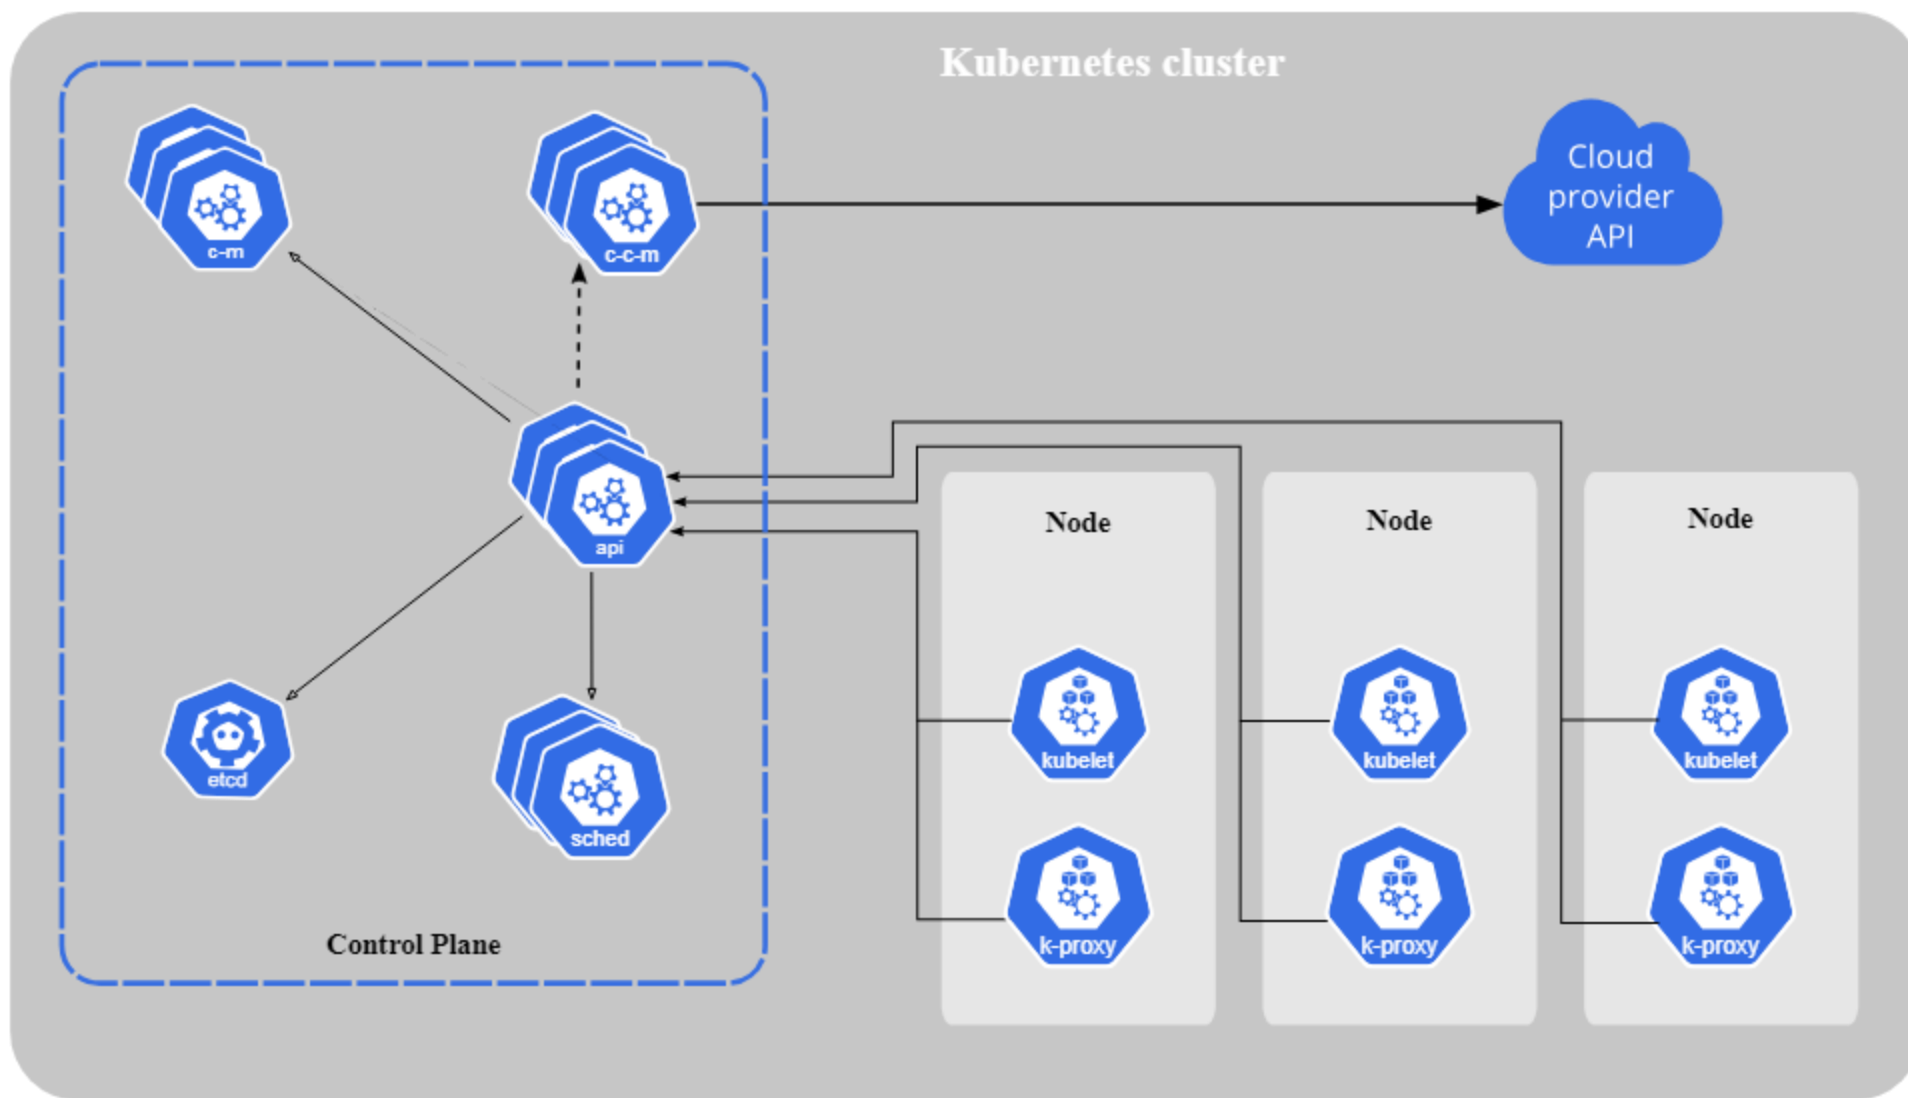
- ▶ Does not provide application-level services, such as
 - ▶ middleware (eg: message buses),
 - ▶ data-processing frameworks (eg: Spark),
 - ▶ databases (eg: MySQL),
 - ▶ caches, nor cluster storage systems (eg: Ceph) as built-in services.
- ▶ Such components
 - ▶ can run on Kubernetes and/or can be accessed by applications running on Kubernetes through portable mechanisms, such as the Open Service Broker.
- ▶ Does not dictate
 - ▶ logging, monitoring, or alerting solutions.
 - ▶ It provides some integrations as proof of concept, and mechanisms to collect and export metrics.
- ▶ Does not provide nor mandate a configuration language/system (eg: Jsonnet).
- ▶ It provides a declarative API that may be targeted by arbitrary forms of declarative specifications.

Kubernetes:

- ▶ Does not provide nor adopt any comprehensive machine configuration, maintenance, management, or self-healing systems.
- ▶ Kubernetes is not a mere orchestration system.
 - ▶ In fact eliminates the need for orchestration.
- ▶ orchestration
 - ▶ execution of a defined workflow: first do A, then B, then C.
- ▶ Kubernetes comprises a set of
 - ▶ independent, composable control processes that continuously drive the current state towards the provided desired state.
 - ▶ It shouldn't matter how we get from A to C.
 - ▶ Centralized control is also not required.
 - ▶ This results in a system that is easier to use and more powerful, robust, resilient, and extensible.

Kubernetes Components

- ▶ When we deploy Kubernetes, we get a cluster.
- ▶ A Kubernetes cluster consists
 - ▶ of a set of worker machines, called nodes, that run containerized applications. Every cluster has at least one worker node.
- ▶ The worker node(s) host the Pods
 - ▶ that are the components of the application workload.
 - ▶ The control plane manages the worker nodes and the Pods in the cluster.
 - ▶ In production environment
 - ▶ the control plane usually runs across multiple computers and
 - ▶ a cluster usually runs multiple nodes, providing fault-tolerance and high availability.



Control Plane Components

- ▶ make global decisions about the cluster eg:
 - ▶ scheduling as well as
 - ▶ detecting and responding to cluster events (eg: starting up a new pod when a deployment's replicas field is unsatisfied).
- ▶ can be run on any machine in the cluster.
- ▶ To keep it simple, we can set up scripts to
 - ▶ start all control plane components on the same machine, and
 - ▶ do not run user containers on this machine.

Control Plane Components : kube-apiserver

- ▶ An API server that exposes the Kubernetes API.
- ▶ The API server is the front end for the Kubernetes control plane.
- ▶ kube-apiserver
 - ▶ Is the main implementation of a Kubernetes API server.
 - ▶ designed to scale horizontally—that is, it scales by deploying more instances.
 - ▶ We can run several instances of kube-apiserver and balance traffic between those instances.

Control Plane Components : etcd

- ▶ Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data.
- ▶ If our Kubernetes cluster uses etcd as its backing store, we need to have a back up plan for the data.

Control Plane Components : kube-scheduler

- ▶ Control plane component that watches for newly created Pods with no assigned node, and selects a node for them to run on.
- ▶ Factors taken into account for scheduling decisions include:
 - ▶ individual and collective resource requirements,
 - ▶ hardware/software/policy constraints,
 - ▶ affinity and anti-affinity specifications,
 - ▶ data locality, inter-workload interference, and deadlines.

Control Plane Components: kube-controller-manager

- ▶ Control plane component that runs controller processes.
- ▶ Logically, each controller is a separate process
 - ▶ to reduce complexity, they are all compiled into a single binary and run in a single process.
- ▶ Many different types of controllers:
 - ▶ Node controller: Responsible for noticing and responding when nodes go down.
 - ▶ Job controller: Watches for Job objects that represent one-off tasks, then creates Pods to run those tasks to completion.
 - ▶ EndpointSlice controller: Populates EndpointSlice objects (to provide a link between Services and Pods).
 - ▶ ServiceAccount controller: Create default ServiceAccounts for new namespaces.

Control Panel Component: cloud-controller-manager

- ▶ embeds cloud-specific control logic.
- ▶ lets us link our cluster into our cloud provider's API, and separates out the components that interact with that cloud platform from components that only interact with our cluster.
- ▶ only runs controllers that are specific to our cloud provider.
- ▶ If we are running Kubernetes on our own premises, or in a learning environment inside our own PC, the cluster does not have a cloud controller manager.
- ▶ Like the kube-controller-manager, it combines several logically independent control loops into a single binary that we run as a single process.
- ▶ We can scale horizontally (run more than one copy) to improve performance or to help tolerate failures.

Control Panel Component: cloud-controller-manager

- ▶ controllers can have cloud provider dependencies:
 - ▶ Node controller:
 - ▶ For checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding
 - ▶ Route controller:
 - ▶ For setting up routes in the underlying cloud infrastructure
 - ▶ Service controller:
 - ▶ For creating, updating and deleting cloud provider load balancers

Node Components

- ▶ run on every node,
- ▶ maintaining running pods and
- ▶ providing the Kubernetes runtime environment.

Node Components: Kubelet

- ▶ An agent that runs on each node in the cluster.
- ▶ makes sure that containers are running in a Pod.
- ▶ takes a set of PodSpecs that are provided through various mechanisms and ensures that the containers described in those PodSpecs are running and healthy.
- ▶ doesn't manage containers which were not created by Kubernetes.

Node Components : kube-proxy

- ▶ A network proxy that runs on each node in our cluster, implementing part of the kubernetes service concept.
- ▶ Maintains network rules on nodes.
 - ▶ These network rules allow network communication to our pods from network sessions inside or outside of your cluster.
- ▶ The operating system packet filtering layer if there is one and it's available.
 - ▶ Otherwise, forwards the traffic itself.

Node Components: Container runtime

- ▶ A fundamental component, empowers Kubernetes to run containers effectively.
- ▶ responsible for managing
 - ▶ the execution and
 - ▶ lifecycle of containers within the Kubernetes environment.
- ▶ Kubernetes supports container runtimes like
 - ▶ containerd,
 - ▶ CRI-O, and any other implementation of the Kubernetes CRI (Container Runtime Interface).

Addons

- ▶ Addons use Kubernetes resources (DaemonSet, Deployment, etc) to implement cluster features.
- ▶ Because these are providing cluster-level features,
 - ▶ namespaced resources for addons belong within the kube-system namespace.

Kubernetes Addon : DNS

- ▶ No addons are strictly required,
- ▶ But all Kubernetes clusters should have cluster DNS, as many examples rely on it.
- ▶ Cluster DNS
 - ▶ is a DNS server, in addition to the other DNS server(s) in our environment, which serves DNS records for Kubernetes services.
- ▶ Containers started by Kubernetes
 - ▶ automatically include this DNS server in their DNS searches.

Kubernetes Addon : Web UI (Dashboard)

- ▶ A general purpose, web-based UI for Kubernetes clusters.
- ▶ allows users to manage and troubleshoot applications running in the cluster, as well as the cluster itself.

Kubernetes Addon : Container Resource Monitoring

- ▶ records generic time-series metrics about containers in a central database,
- ▶ and provides a UI for browsing that data.

Kubernetes Addon : Cluster-level Logging

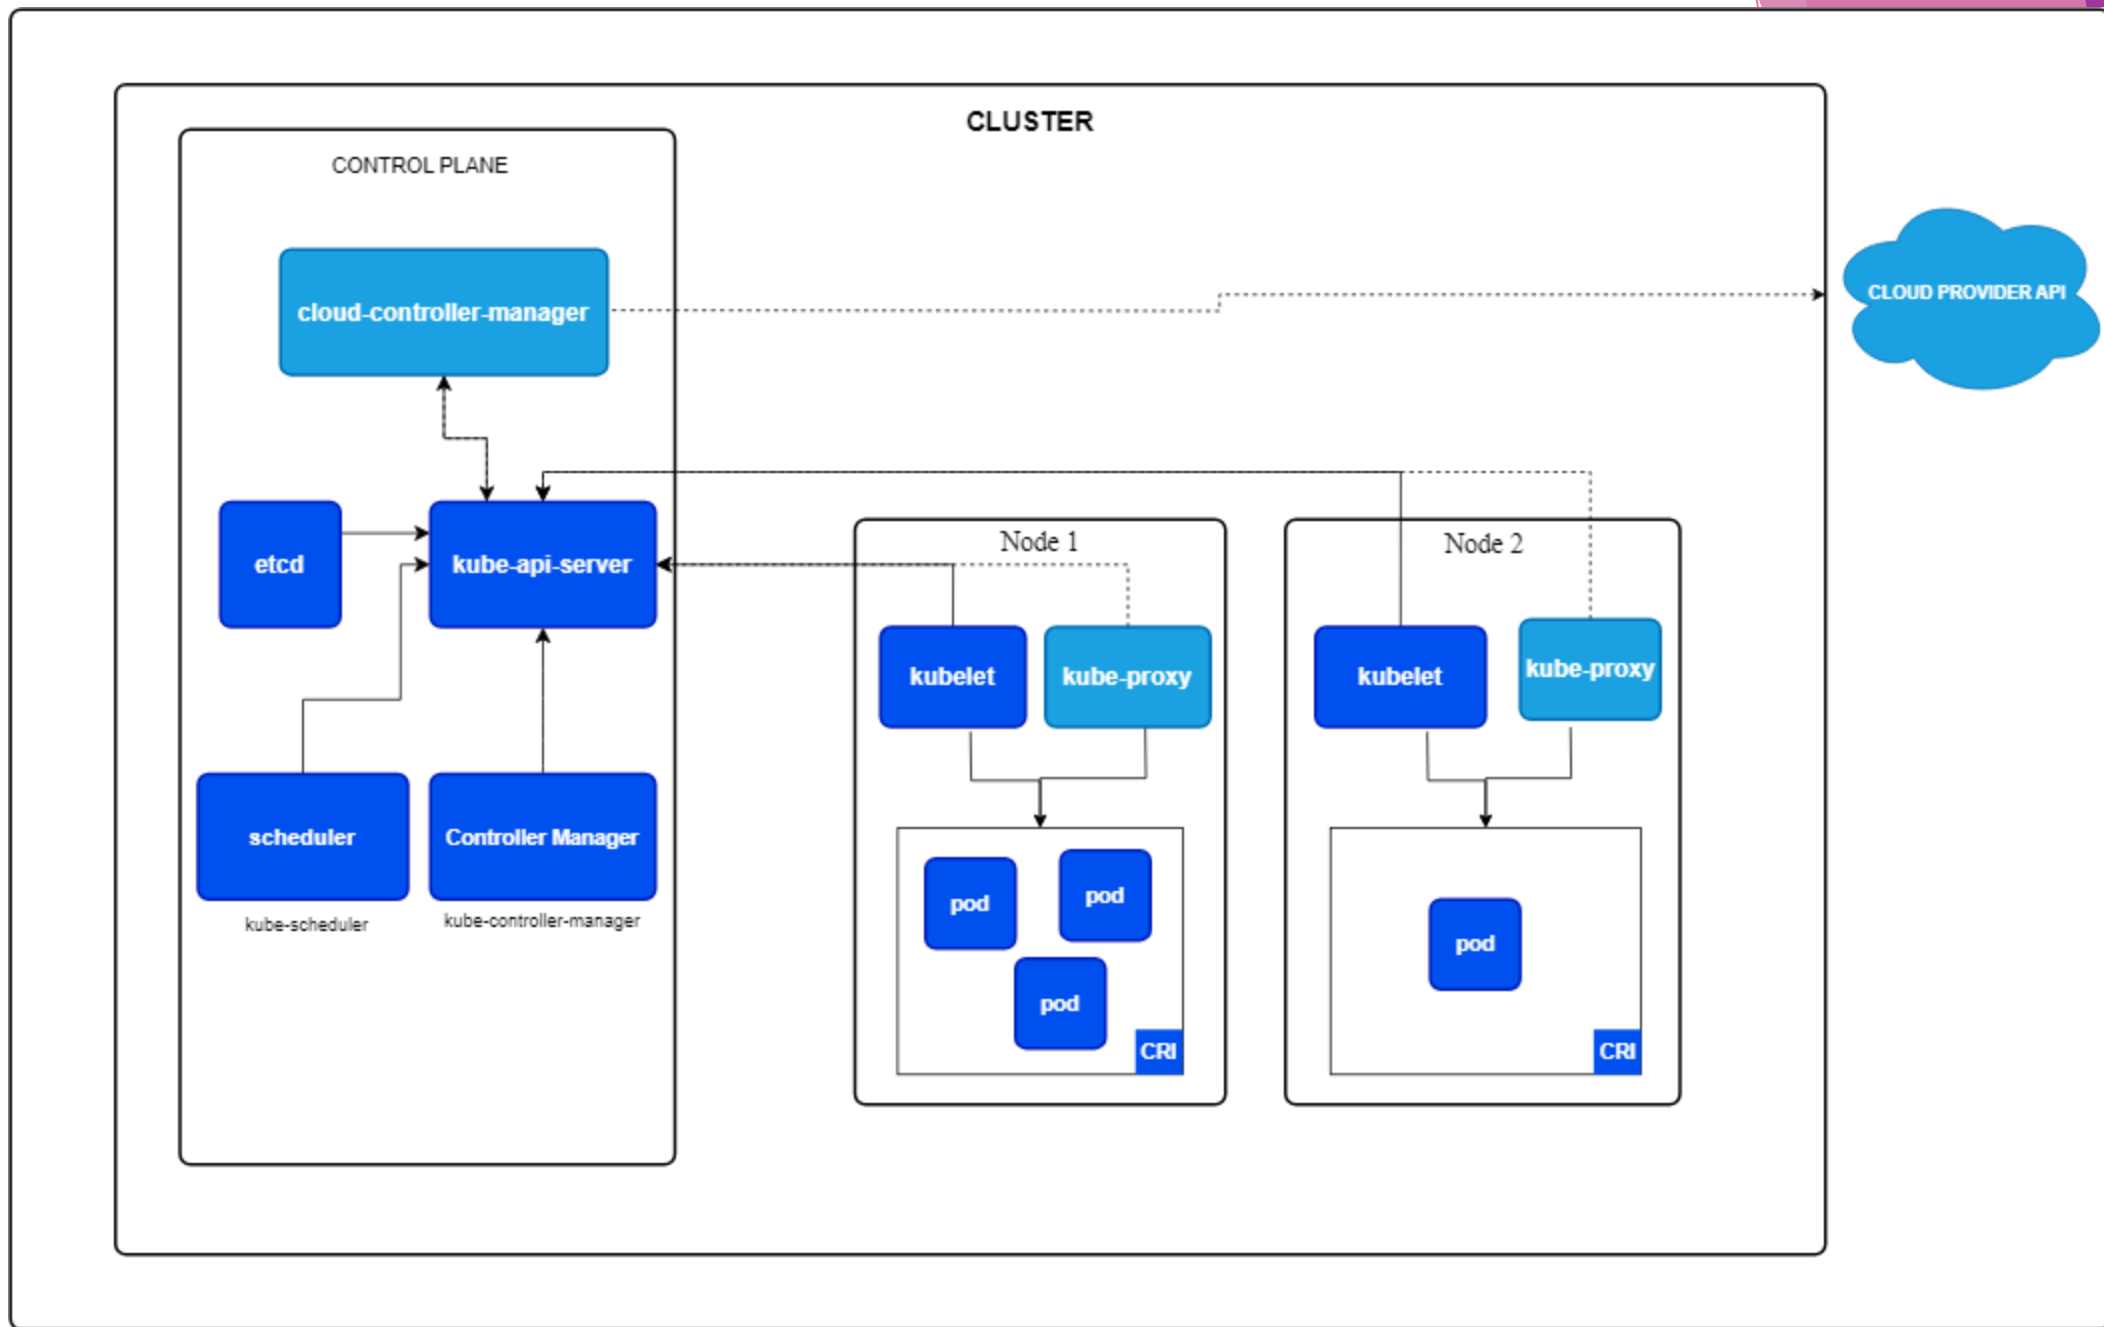
- ▶ A cluster-level logging mechanism
- ▶ responsible for saving container logs to a central log store with search/browsing interface.

Kubernetes Addon : Network Plugins

- ▶ software components that implement the container network interface (CNI) specification.
- ▶ responsible for allocating IP addresses to pods and enabling them to communicate with each other within the cluster.

Any Questions?

Kubernetes Cluster Architecture



Use Case

- ▶ Create a deployment to rollout a replica set
- ▶ Declare a new state for a pod
- ▶ Rollback to an earlier deployment revision
- ▶ Scale up the Deployment to facilitate more load
- ▶ Pause the rollout of the deployment
- ▶ Use the status of the deployment
- ▶ Clean up of replicaset not needed anymore